

## 1. Introduction

In this project we develop and evaluate a convolutional neural network—**ColabOptimizedCNN**—to classify plant leaf images into multiple disease categories. We implemented data loading, model training with early stopping, advanced evaluation (precision-recall, ROC, confusion matrix), feature-space visualization (PCA, t-SNE, UMAP), and model complexity analysis. The following document breaks down each component in simple, detailed steps, then interprets the results shown in the attached charts.

## 2. Environment Setup and Data Preparation

- Package installation:** Torch, torchvision, Matplotlib, scikit-learn, UMAP, etc.
- Mount Google Drive:** so we can load and save large datasets and models directly.
- Zip extraction:** we locate and unzip `Plant_leaf_diseases_dataset_with_augmentation.zip` into `/content/extracted_dataset`.
- Automatic root detection:** `find_dataset_root(...)` scans subfolders until it finds class directories containing images.
- Dataset exploration:** prints the first 10 classes and their image counts, plus the total number of images and classes.
- Train/Validation/Test split:** 70% train, 15% validation, 15% test, with a fixed random seed for reproducibility.

## 3. Image Transforms and DataLoaders

- Dynamic sizing:** If a GPU is available, images are resized to 128×128 then center-cropped to 112×112; on CPU we use 96→84 to speed up training.
- Augmentations** (training only): random horizontal flips, rotations, brightness/contrast jitter to improve generalization.
- Normalization:** standard ImageNet means and standard deviations.
- DataLoaders:** batch size 64 on GPU (32 on CPU), two workers, pinned memory if possible, and `SubsetRandomSampler` for each split.

## 4. Model Architecture: ColabOptimizedCNN

A lightweight, mobile-friendly CNN inspired by depthwise separable convolutions:

Stage	Layers	Purpose
1	Conv2d(3→32), BatchNorm, ReLU6	Initial feature extraction
2	<code>_make_depthwise_block(32→64, stride=1)</code>	Depthwise + pointwise conv block
3	<code>_make_depthwise_block(64→128, stride=2)</code>	Downsampling block
...	Additional depthwise blocks up to 512 channels	Progressive deepening
Final	AdaptiveAvgPool → Flatten → Dropout → Linear(512→256) → ReLU → Dropout → Linear(256→num_classes)	Classification head

- Depthwise separable blocks** split spatial and channel mixing, reducing parameters.
- Initialization:** Kaiming for conv layers, constant for batch norm, small normal for linear layers.
- Feature hook:** `get_features()` returns the 512-dim vector before the classifier for later clustering.

Total parameters: ~410 k (≈1.57 MB) Trainable parameters: 410 k

## 5. Training Procedure

- Optimizer:** AdamW with weight decay 1e-4, initial LR 0.001.
- Scheduler:** StepLR reduces LR by ½ every 7 epochs.
- Early stopping:** stops if validation accuracy does not improve for 5 consecutive epochs.
- Logging:** per-batch loss every 50 batches, plus per-epoch summary:
  - Train loss & accuracy
  - Val loss & accuracy
  - Learning rate
  - Early stopping counter

Training run: 20 epochs, best validation accuracy ≈ 0.9882 at epoch 20.

---

## 6. Test-set Evaluation Metrics

---

After loading the best model checkpoint, we compute:

- **Loss** on test set: ~0.0364
- **Accuracy**: 0.9873
- **Macro Precision**: 0.9849
- **Macro Recall**: 0.9839
- **Macro F1-score**: 0.9842
- **Micro Precision**: 0.9873
- **Micro Recall**: 0.9873
- **Micro F1-score**: 0.9873

These high scores reflect excellent overall and per-class performance.

---

## 7. Precision–Recall and ROC Curves

---

### Precision–Recall (first image set)

- **Micro-averaged**: nearly perfect curve hugging the top-right corner (precision  $\approx 1.0$  for recall up to 0.99).
- **Macro-averaged**: similarly high, indicating all classes perform well.
- **Per-class (top 5)**: each of the first five classes shows near-unit precision and recall.

### ROC (second image set)

- **Micro AUC**: 1.000
- **Macro AUC**:  $\approx 0.9999$
- **Per-class AUC (top 5)**: all AUC = 1.000.

**Interpretation**: The model separates positive vs. negative examples perfectly across classes, with negligible overlap.

---

## 8. Feature-Space Clustering (third image set)

---

We extract 512-dim feature vectors for 1,000 test samples, then visualize with:

1. **PCA**: first two components capture ~28.1% variance. Clusters start to form but many overlap.
2. **t-SNE**: tighter, well-separated clusters showing that learned features discriminate classes.
3. **UMAP**: similar separation, sometimes revealing substructure within classes.
4. **Class distribution** bar chart: shows sample counts per class in this subset (some classes more frequent than others).

**Takeaway**: feature extractor produces linearly separable clusters for most classes, confirming strong representation learning.

---

## 9. Model Complexity Analysis

---

From the printed layer-by-layer breakdown:

- **Conv2d layers**: ~64.5% of total parameters
- **BatchNorm2d layers**: ~1.1%
- **Linear layers**: ~34.4%

Other plots:

- **Final layers bar chart**: shows that the two final linear transforms (256→num\_classes) contain the majority of parameters.
- **Memory usage**: parameters and gradients each  $\approx 1.56$  MB, activations estimate  $\approx 2.57$  MB per batch.
- **Trainable vs. non-trainable**: all parameters are trainable.

**Conclusion**: At ~410 k parameters and <2 MB size, this model is compact and suitable for resource-limited environments.

---

## 10. Training Dynamics and Overfitting Check (last image set)

---

1. **Loss curves**: training and validation loss both decrease smoothly, with validation always below training after epoch 2, indicating no overfitting.
  2. **Validation accuracy**: climbs from ~0.73 to ~0.99, plateauing near the best value.
  3. **Learning rate schedule**: starts at 1e-3, halves at epoch 7 to 5e-4, then halves again at epoch 14 to 2.5e-4.
  4. **Overfitting indicator** (val – train loss): stays negative or near zero, confirming good generalization.
  5. **Test-set metrics bar chart**: re-plots accuracy, precision, recall, F1 for macro and micro, all near 0.99.
  6. **Convergence analysis**: moving average of val-accuracy shows stable convergence by epoch 10.
  7. **Training summary box**: lists final epochs, best val acc, test acc, total params, size, final LR, and approximate training time (~40 min).
-

## 11. Sample Predictions

---

The final section (not plotted as a chart) prints 5 random test images with:

- **True vs. predicted class**
- **Confidence score** of top prediction
- **Top 3 predicted classes** with probabilities

All sample predictions show correct labels with high confidence ( $\approx 1.000$ ), further illustrating the model's reliability.

---

## 12. Conclusion

---

- **Performance:** 98.7% test accuracy with balanced macro/micro metrics.
- **Efficiency:** <1 MB model, fast inference, suitable for deployment on mobile/edge.
- **Interpretability:** clear feature clustering, simple architecture, and transparent complexity breakdown.

**Future work** could explore real-time deployment on smartphones, continued data augmentation for robustness, or lighter backbones for even smaller footprints.

---

*Prepared by Group 13 – BSc Computer Science, June 2025*